GraphSlim on ogbn-arxiv

169,343 nodes · full undirected graph · one RTX PRO 6000 Blackwell

Lower is better ↓

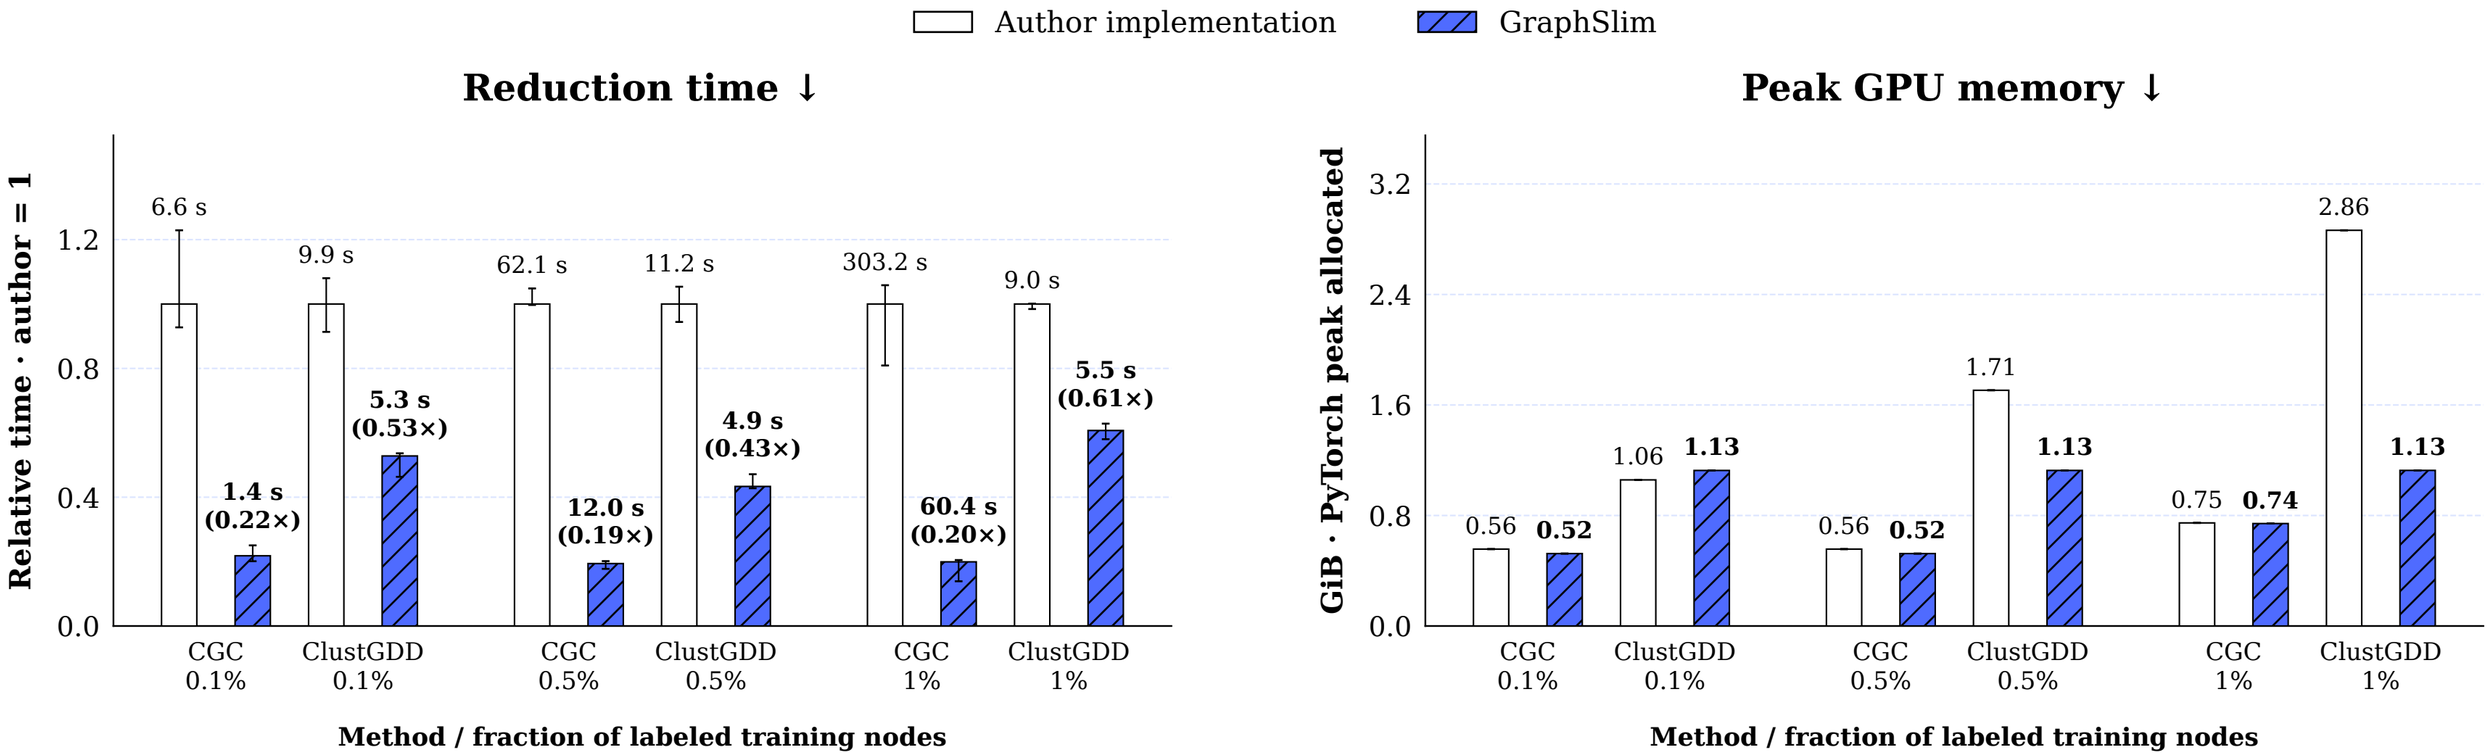

3 reductions per bar; whiskers: observed min–max. Time is relative to each author's median (1); lower is better for both metrics.

0.1%, 0.5%, 1% of training nodes retain 90, 454, 909 nodes. Timing and GPU peaks exclude downstream GCN evaluation.