GraphSlim vs. author implementations

Matched inputs and output sizes · same device within each pair · all six comparisons

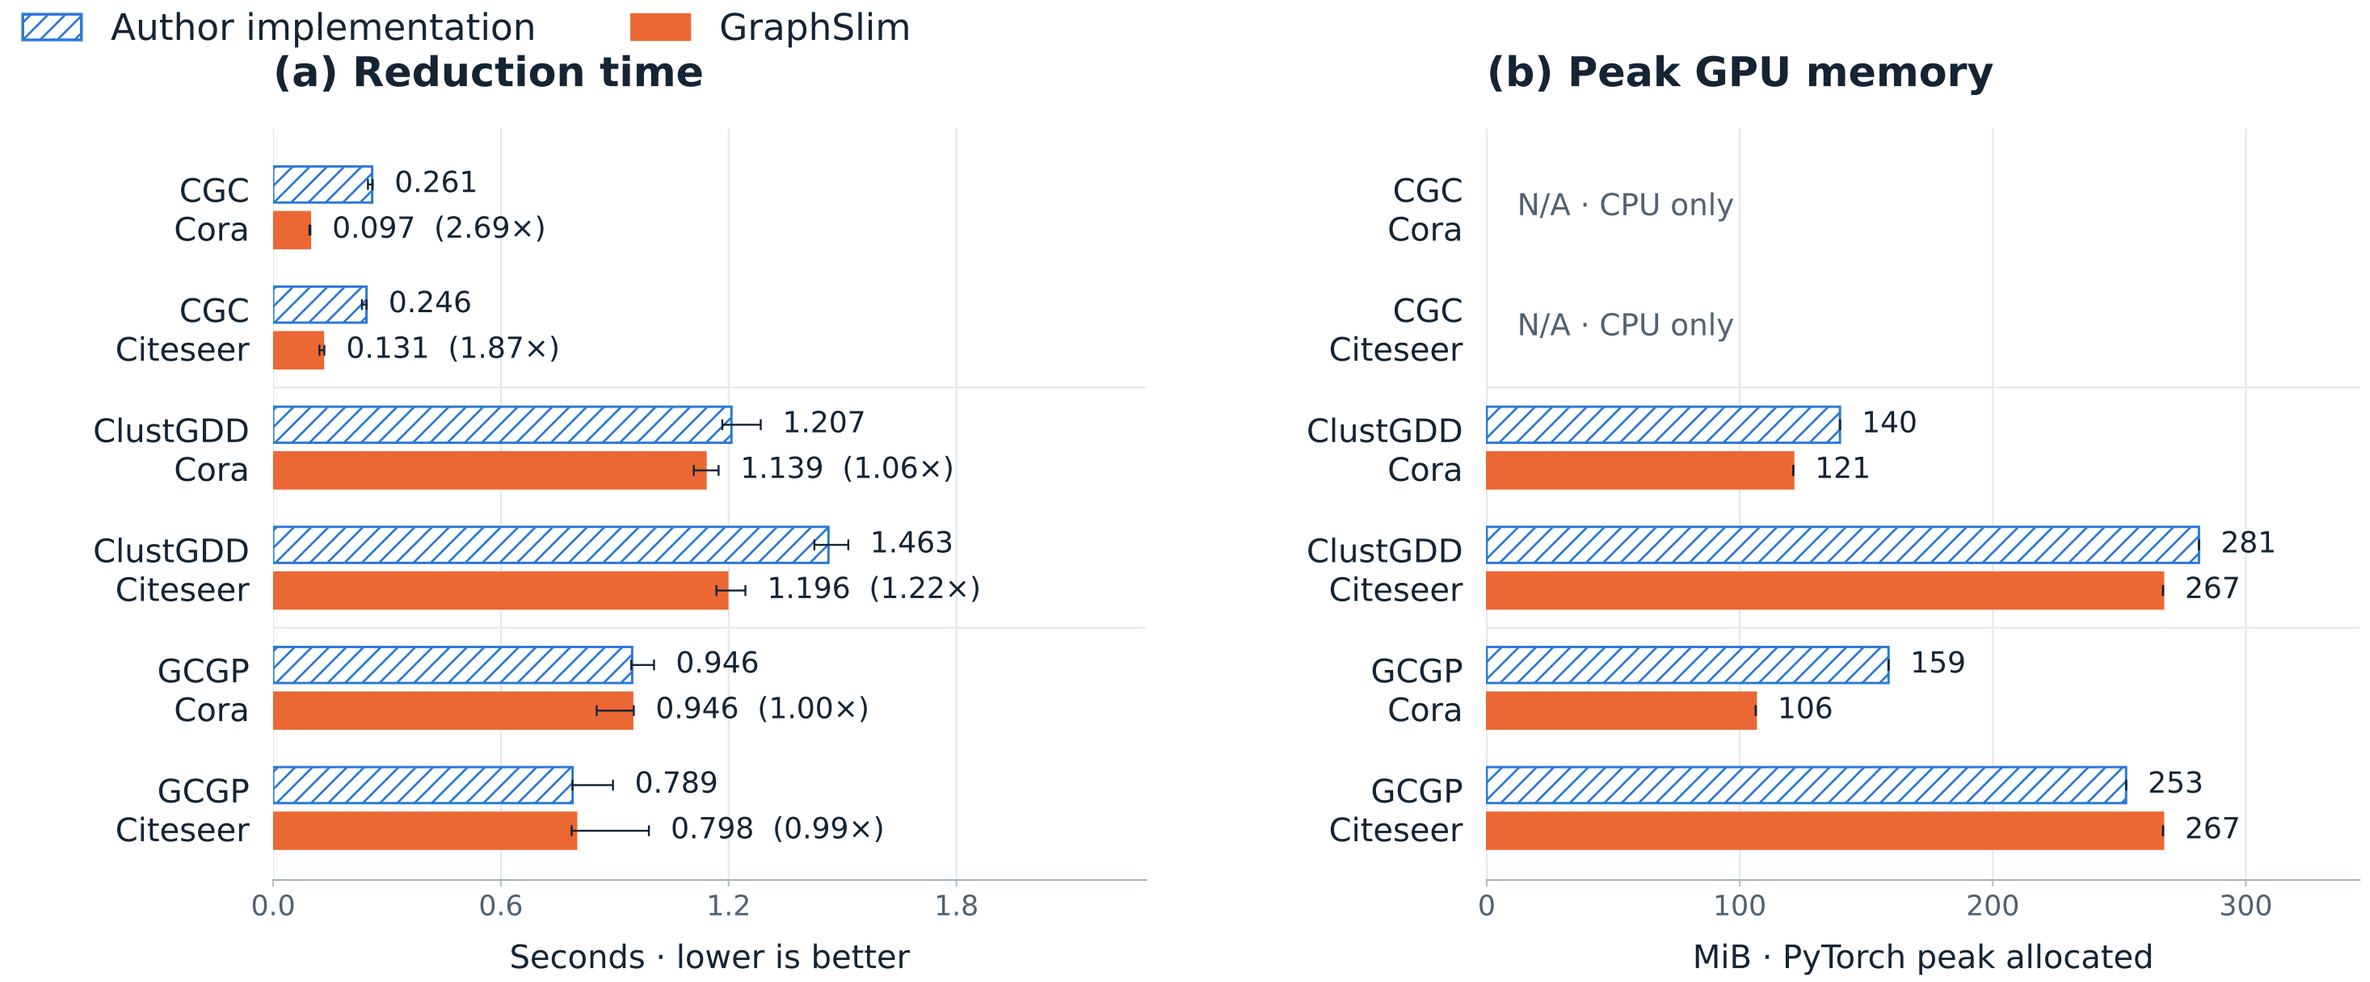

Time / memory: medians of 3 reductions; whiskers: observed min–max. All pairs retain 50% of labeled training nodes.

CGC: 4 CPU threads (GPU memory N/A). ClustGDD / GCGP: one RTX PRO 6000 Blackwell. GPU memory: PyTorch peak allocated.