GraphSlim vs. author implementations

Matched inputs and output sizes · same device within each pair · all six comparisons

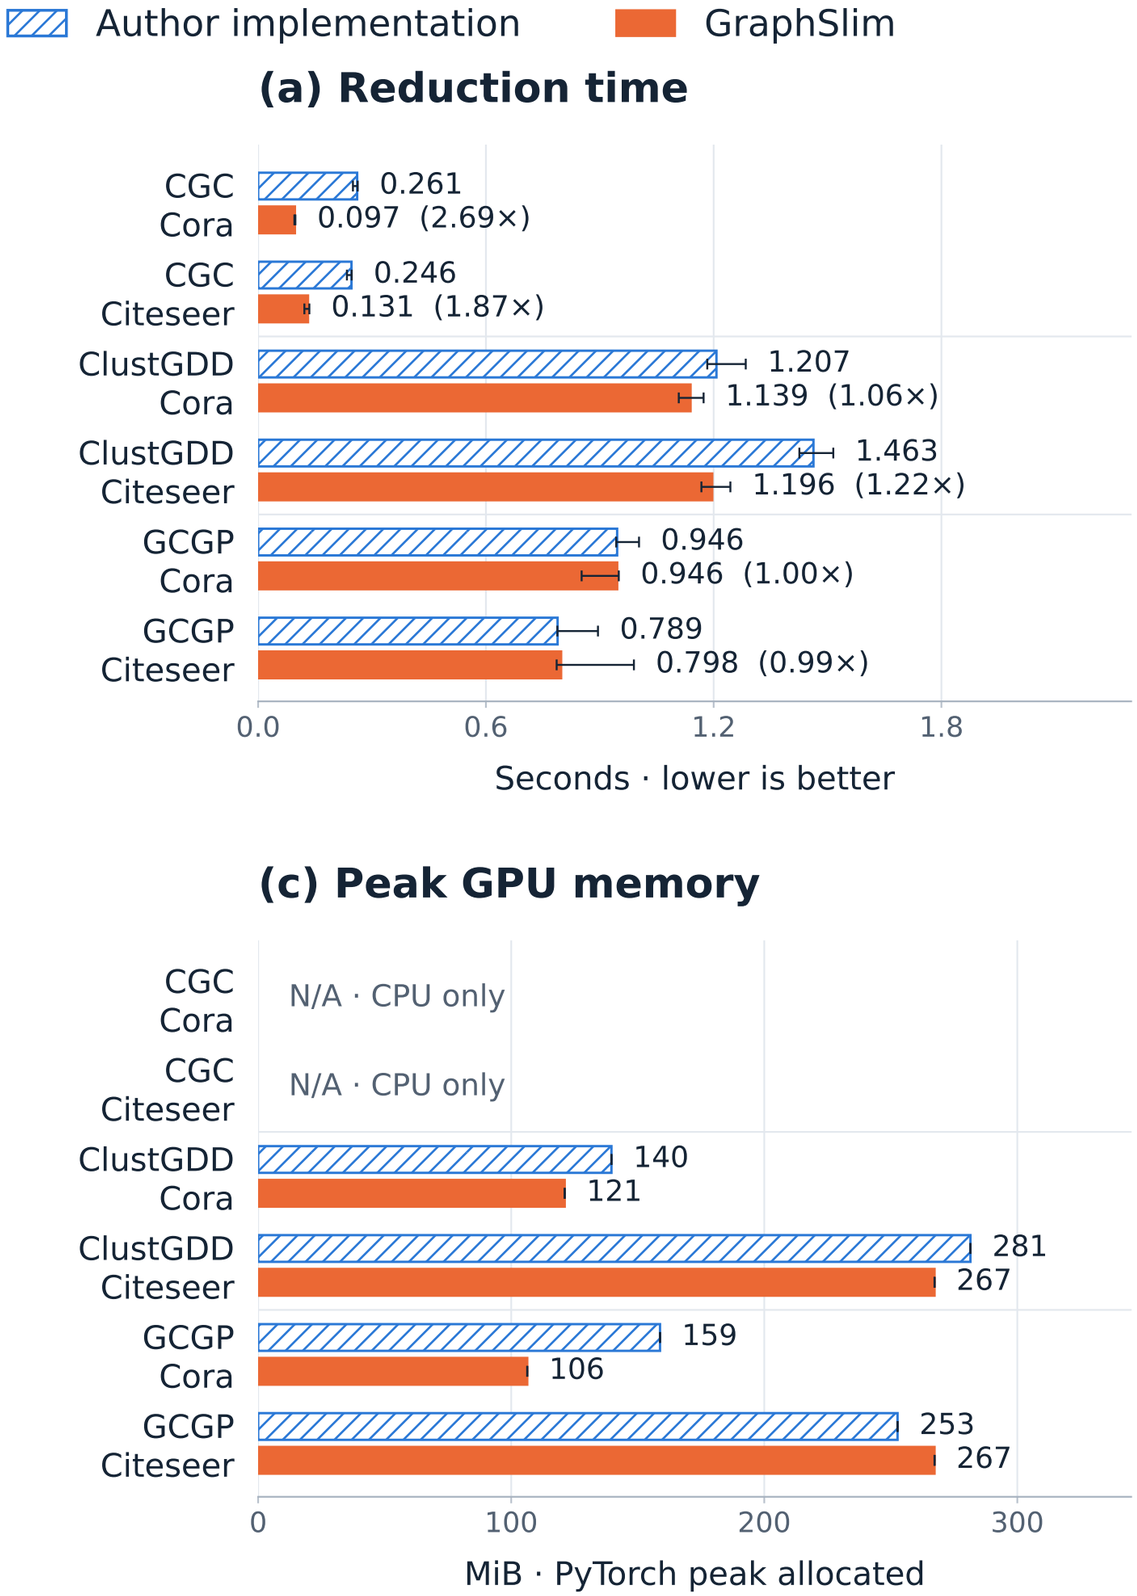

Time / memory: medians of 3 reductions; whiskers: observed min–max. Accuracy: gap between means (GCN: 30 fits; GP: 3 graphs).

CGC: 4 CPU threads. ClustGDD / GCGP: one RTX PRO 6000 Blackwell. RSS sampled every 5 ms; CUDA is allocated, not reserved memory.